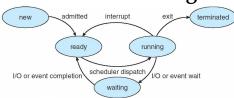
CPU Scheduling P1 P2 ... PN

- The scheduling problem:
 - Have K jobs ready to run
 - Have $N \ge 1$ CPUs
 - Which jobs to assign to which CPU(s)
- When do we make decision?

Scheduling criteria

- Why do we care?
 - What goals should we have for a scheduling algorithm?

CPU Scheduling



- Scheduling decisions may take place when a process:
 - 1. Switches from running to waiting state
 - 2. Switches from running to ready state
 - 3. Switches from waiting to ready
 - 4. Exits
- Non-preemptive schedules use 1 & 4 only
- Preemptive schedulers run at all four points

Scheduling criteria

- Why do we care?
 - What goals should we have for a scheduling algorithm?
- Throughput # of procs that complete per unit time
 - Higher is better
- Turnaround time time for each proc to complete
 - Lower is better
- Response time time from request to first response (e.g., key press to character echo, not launch to exit)
 - Lower is better
- Above criteria are affected by secondary criteria
 - CPU utilization keep the CPU as busy as possible
 - Waiting time time each proc waits in ready queue

Example: FCFS Scheduling

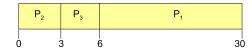
- Run jobs in order that they arrive
 - Called "First-come first-served" (FCFS)
 - E.g., Say P_1 needs 24 sec, while P_2 and P_3 need 3.
 - Say P_2 , P_3 arrived immediately after P_1 , get:



- Dirt simple to implement—how good is it?
- Throughput: 3 jobs / 30 sec = 0.1 jobs/sec
- Turnaround Time: $P_1: 24, P_2: 27, P_3: 30$
 - Average TT: (24 + 27 + 30)/3 = 27
- Can we do better?

FCFS continued

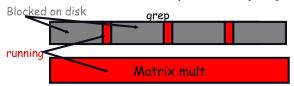
- Suppose we scheduled P_2 , P_3 , then P_1
 - Would get:



- Throughput: 3 jobs / 30 sec = 0.1 jobs/sec
- Turnaround time: $P_1 : 30, P_2 : 3, P_3 : 6$
 - Average TT: (30 + 3 + 6)/3 = 13 much less than 27
- Lesson: scheduling algorithm can reduce TT
 - Minimize waiting time to minimize TT
- What about throughput?

I/O devices just special CPUs

- An I/O device is like a special purpose CPU
 - "special purpose" = disk drive can only run a disk job, tape drive a tape job, . . .
- Implication: 1-CPU system with n I/O devices is like an n + 1-CPU multiprocessor
 - Result: all I/O devices + CPU busy ⇒ n+1 fold speedup!



- Overlap them just right? throughput will be doubled

Bursts of computation & I/O

Jobs contain I/O and computation

- Bursts of computation
- Then must wait for I/O

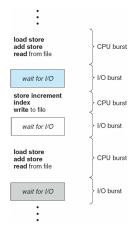
• To Maximize throughput

- Must maximize CPU utilization
- Also maximize I/O device utilization

· How to do?

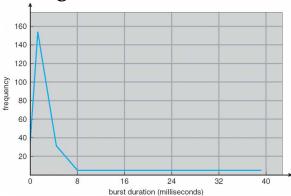
6/32

- Overlap I/O & computation from multiple jobs
- Means response time very important for I/O-intensive jobs: I/O device will be idle until job gets small amount of CPU to issue next I/O request



7/32

Histogram of CPU-burst times



• What does this mean for FCFS?

FCFS Convoy effect

- CPU bound jobs will hold CPU until exit or I/O (but I/O rare for CPU-bound thread)
 - long periods where no I/O requests issued, and CPU held
 - Result: poor I/O device utilization
- Example: one CPU-bound job, many I/O bound
 - CPU bound runs (I/O devices idle)
 - CPU bound blocks
 - I/O bound job(s) run, quickly block on I/O
 - CPU bound runs again
 - I/O completes
 - CPU bound still runs while I/O devices idle (continues?)
- Simple hack: run process whose I/O completed?
 - What is a potential problem?

SJF Scheduling

- Shortest-job first (SJF) attempts to minimize TT
- Two schemes:
 - nonpreemptive once CPU given to the process it cannot be preempted until completes its CPU burst
 - preemptive if a new process arrives with CPU burst length less than remaining time of current executing process, preempt (Know as the Shortest-Remaining-Time-First or SRTF)
- What does SJF optimize?

SJF Scheduling

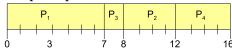
- Shortest-job first (SJF) attempts to minimize TT
- Two schemes:
 - nonpreemptive once CPU given to the process it cannot be preempted until completes its CPU burst
 - preemptive if a new process arrives with CPU burst length less than remaining time of current executing process, preempt (Know as the Shortest-Remaining-Time-First or SRTF)
- What does SJF optimize?
 - gives minimum average waiting time for a given set of processes

732

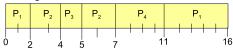
Examples

Process	Arrival Time	Burst Time
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

• Non-preemptive



• Preemptive



• Drawbacks?

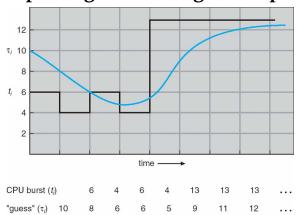
SJF limitations

• Doesn't always minimize average turnaround time

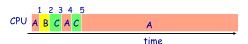
- Only minimizes waiting time
- Example where turnaround time might be suboptimal?
- Can lead to unfairness or starvation
- In practice, can't actually predict the future
- But can estimate CPU burst length based on past
 - Exponentially weighted average a good idea
 - t_n actual length of proc's n^{th} CPU burst
 - τ_{n+1} estimated length of proc's $n+1^{\rm st}$
 - Choose parameter α where $0 < \alpha \le 1$
 - Let $\tau_{n+1} = \alpha t_n + (1 \alpha)\tau_n$

12/32

Exp. weighted average example



Round robin (RR) scheduling



• Solution to fairness and starvation

- Preempt job after some time slice or quantum
- When preempted, move to back of FIFO queue
- (Most systems do some flavor of this)

• Advantages:

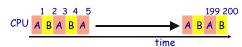
11/32

- Fair allocation of CPU across jobs
- Low average waiting time when job lengths vary
- Good for responsiveness if small number of jobs
- · Disadvantages?

13/32

RR disadvantages

- Varying sized jobs are good...
- but what about same-sized jobs?
- Assume 2 jobs of time=100 each:



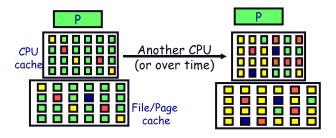
- What is average completion time?
- How does that compare to FCFS?

Context switch costs

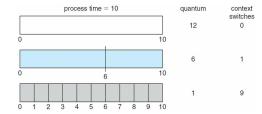
• What is the cost of a context switch?

Context switch costs

- What is the cost of a context switch?
- Brute CPU time cost in kernel
 - Save and restore resisters, etc.
 - Switch address spaces (expensive instructions)
- Indirect costs: cache, buffer cache, & TLB misses



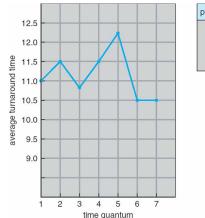
Time quantum



- How to pick quantum?
 - Want much larger than context switch cost
 - But not so large system reverts to FCFS
- Typical values: 10-100 msec

17/22

Turnaround time vs. quantum



process	time
P ₁	6
P_2	3
P_3	1
P_4	7

Two-level scheduling

- Switching to swapped out process very expensive
 - Swapped out process has most pages on disk
 - Will have to fault them all in while running
 - One disk access costs 10ms. On 1GHz machine, 10ms = 10 million cycles!
- Context-switch-cost aware scheduling
 - Run in core subset for "a while"
 - Then move some between disk and memory
 - How to pick subset? Hot to define "a while"?

,

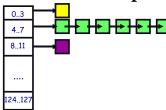
Priority scheduling

- A priority number (integer) is associated with each process
 - E.g., smaller priority number means higher priority
- Give CPU to the process with highest priority
 - Can be done preemptively or non-preemptively
- Note SJF is a priority scheduling where priority is the predicted next CPU burst time
- Starvation low priority processes may never execute
- Solution?

Priority scheduling

- A priority number (integer) is associated with each process
 - E.g., smaller priority number means higher priority
- Give CPU to the process with highest priority
 - Can be done preemptively or non-preemptively
- Note SJF is a priority scheduling where priority is the predicted next CPU burst time
- Starvation low priority processes may never execute
- Solution?
 - Aging increase a process's priority as it waits

Multilevel feeedback queues (BSD)



- Every runnable proc. on one of 32 run queues
 - Kernel runs proc. on highest-priority non-empty queue
 - Round-robins among processes on same queue
- Process priorities dynamically computed
 - Processes moved between queues to reflect priority changes
 - If a proc. gets higher priority than running proc., run it
- Idea: Favor interactive jobs that use less CPU

Process priority

- p_nice user-settable weighting factor
- p_estcpu per-process estimated CPU usage
 - Incremented whenever timer interrupt found proc. running
 - Decayed every second while process runnable

$$\texttt{p_estcpu} \leftarrow \left(\frac{2 \cdot load}{2 \cdot load + 1}\right) \texttt{p_estcpu} + \texttt{p_nice}$$

• Run queue determined by p_usrpri/4

$$\texttt{p_usrpri} \leftarrow 50 + \left(\frac{\texttt{p_estcpu}}{4}\right) + 2 \cdot \texttt{p_nice}$$

(value clipped if over 127)

21/32

22/32

Sleeping process increases priority

- p_estcpu not updated while asleep
 - Instead p_slptime keeps count of sleep time
- When process becomes runnable

$$p_{\texttt{estcpu}} \leftarrow \left(\frac{2 \cdot load}{2 \cdot load + 1}\right)^{p_{\texttt{eslptime}}} \times p_{\texttt{estcpu}}$$

- Approximates decay ignoring nice and past loads

Limitations of BSD scheduler

- Hard to have isolation / prevent interference
 - Priorities are absolute
- Can't transfer priority (e.g., to server on RPC)
- No flexible control
 - E.g., In monte carlo simulations, error is 1/sqrt(N) after N trials
 - Want to get quick estimate from new computation
 - Leave a bunch running for a while to get more accurate results
- Multimedia applications
 - Often fall back to degraded quality levels depending on resources
 - Want to control quality of different streams

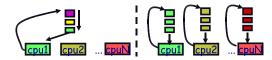
24/3

Real-time scheduling

- Two categories:
 - Soft real time—miss deadline and CD will sound funny
 - Hard real time—miss deadline and plane will crash
- System must handle periodic and aperiodic events
 - E.g., procs A, B, C must be scheduled every 100, 200, 500 msec, require 50, 30, 100 msec respectively
 - Schedulable if $\sum \frac{CPU}{\text{period}} \le 1$ (not counting switch time)
- Variety of scheduling strategies
 - E.g., first deadline first (works if schedulable)

Multiprocessor scheduling issues

- Must decide more than which process to run
 - Must decide on which CPU to run it
- Moving between CPUs has costs
 - More cache misses, depending on arch more TLB misses too
- Affinity scheduling—try to keep threads on same CPU

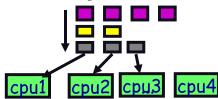


- But also prevent load imbalances
- Do cost-benefit analysis when deciding to migrate

32

Multiprocessor scheduling (cont)

- Want related processes scheduled together
 - Good if threads access same resources (e.g., cached files)
 - Even more important if threads communicate often, otherwise must context switch to communicate
- Gang scheduling—schedule all CPUs synchronously
 - With synchronized quanta, easier to schedule related processes/threads together



27/32

Thread scheduling

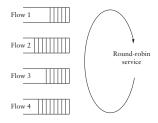
- With thread library, have two scheduling decisions:
 - Local Scheduling Threads library decides which user thread to put onto an available kernel thread
 - Global Scheduling Kernel decides which kernel thread to run next
- Can expose to the user
 - E.g., pthread_attr_setscope allows two choices
 - PTHREAD_SCOPE_SYSTEM thread scheduled like a process (effectively one kernel thread bound to user thread
 - Will return ENOTSUP in user-level pthreads implementation)
 - PTHREAD_SCOPE_PROCESS thread scheduled within the current process (may have multiple user threads multiplexed onto kernel threads)

Thread dependencies

- Priority inversion e.g., T_1 at high priority, T_2 at low
 - T₂ acquires lock L.
 - Scene 1: T_1 tries to acquire L, fails, spins. T_2 never gets to run.
 - Scene 2: T₁ tries to acquire L, fails, blocks. T₃ enters system at medium priority. T₂ never gets to run.
- Scheduling = deciding who should make progress
 - Obvious: a thread's importance should increase with the importance of those that depend on it.
 - Naïve priority schemes violate this
- "Priority donation"
 - Thread's priority scales w. priority of dependent threads

Fair Queuing (FQ)

- Digression: packet scheduling problem
 - Which network packet to send next over a link?
 - Problem inspired some algorithms we will see next time
- For ideal fairness, would send one bit from each flow
 - In weighted fair queuing (WFQ), more bits from some flows



• Complication: must send whole packets

FQ Algorithm

- Suppose clock ticks each time a bit is transmitted
- Let P_i denote the length of packet i
- Let S_i denote the time when start to transmit packet i
- Let F_i denote the time when finish transmitting packet i
- $F_i = S_i + P_i$
- When does router start transmitting packet *i*?
 - If arrived before router finished packet i-1 from this flow, then immediately after last bit of i-1 (F_{i-1})
 - If no current packets for this flow, then start transmitting when arrives (call this A_i)
- Thus: $F_i = \max(F_{i-1}, A_i) + P_i$

FQ Algorithm (cont)

- For multiple flows
 - Calculate F_i for each packet that arrives on each flow
 - Treat all F_i s as timestamps
 - Next packet to transmit is one with lowest timestamp
- Not perfect: can't preempt current packet
- Example:

